



jerrybai1995 Added missing files; updated documentations

Latest commit 029afeb on Jun 6, 2019

History

2 contributors

146 lines (126 sloc) | 6.48 KB

Raw

Blame



```
1 import torch
2 from torch import nn
3 import torch.nn.functional as F
4
5 from modules.transformer import TransformerEncoder
6
7
8 class MULTModel(nn.Module):
9     def __init__(self, hyp_params):
10         """
11         Construct a Mult model.
12         """
13         super(MULTModel, self).__init__()
14         self.orig_d_l, self.orig_d_a, self.orig_d_v = hyp_params.orig_d_l, hyp_params.orig_d_a, hyp_params.orig_d_v
15         self.d_l, self.d_a, self.d_v = 30, 30, 30
16         self.vonly = hyp_params.vonly
17         self.aonly = hyp_params.aonly
18         self.lonly = hyp_params.lonly
19         self.num_heads = hyp_params.num_heads
20         self.layers = hyp_params.layers
21         self.attn_dropout = hyp_params.attn_dropout
22         self.attn_dropout_a = hyp_params.attn_dropout_a
23         self.attn_dropout_v = hyp_params.attn_dropout_v
24         self.relu_dropout = hyp_params.relu_dropout
25         self.res_dropout = hyp_params.res_dropout
26         self.out_dropout = hyp_params.out_dropout
27         self.embed_dropout = hyp_params.embed_dropout
28         self.attn_mask = hyp_params.attn_mask
29
30         combined_dim = self.d_l + self.d_a + self.d_v
31
32         self.partial_mode = self.lonly + self.aonly + self.vonly
33         if self.partial_mode == 1:
34             combined_dim = 2 * self.d_l # assuming d_l == d_a == d_v
35         else:
36             combined_dim = 2 * (self.d_l + self.d_a + self.d_v)
37
38         output_dim = hyp_params.output_dim # This is actually not a hyperparameter :-)
39
40         # 1. Temporal convolutional layers
41         self.proj_l = nn.Conv1d(self.orig_d_l, self.d_l, kernel_size=1, padding=0, bias=False)
42         self.proj_a = nn.Conv1d(self.orig_d_a, self.d_a, kernel_size=1, padding=0, bias=False)
43         self.proj_v = nn.Conv1d(self.orig_d_v, self.d_v, kernel_size=1, padding=0, bias=False)
44
45         # 2. Crossmodal Attentions
46         if self.lonly:
47             self.trans_l_with_a = self.get_network(self_type='la')
48             self.trans_l_with_v = self.get_network(self_type='lv')
49         if self.aonly:
50             self.trans_a_with_l = self.get_network(self_type='al')
51             self.trans_a_with_v = self.get_network(self_type='av')
52         if self.vonly:
53             self.trans_v_with_l = self.get_network(self_type='vl')
54             self.trans_v_with_a = self.get_network(self_type='va')
55
56         # 3. Self Attentions (Could be replaced by LSTMs, GRUs, etc.)
57         # [e.g., self.trans_x_mem = nn.LSTM(self.d_x, self.d_x, 1)
58         self.trans_l_mem = self.get_network(self_type='l_mem', layers=3)
59         self.trans_a_mem = self.get_network(self_type='a_mem', layers=3)
60         self.trans_v_mem = self.get_network(self_type='v_mem', layers=3)
61
62         # Projection layers
63         self.proj1 = nn.Linear(combined_dim, combined_dim)
64         self.proj2 = nn.Linear(combined_dim, combined_dim)
65         self.out_layer = nn.Linear(combined_dim, output_dim)
66
67     def get_network(self, self_type='l', layers=-1):
68         if self_type in ['l', 'al', 'vl']:
69             embed_dim, attn_dropout = self.d_l, self.attn_dropout
70         elif self_type in ['a', 'la', 'va']:
71             embed_dim, attn_dropout = self.d_a, self.attn_dropout_a
72         elif self_type in ['v', 'lv', 'av']:
73             embed_dim, attn_dropout = self.d_v, self.attn_dropout_v
74         elif self_type == 'l_mem':
75             embed_dim, attn_dropout = 2*self.d_l, self.attn_dropout
76         elif self_type == 'a_mem':
77             embed_dim, attn_dropout = 2*self.d_a, self.attn_dropout
78         elif self_type == 'v_mem':
79             embed_dim, attn_dropout = 2*self.d_v, self.attn_dropout
80         else:
81             raise ValueError("Unknown network type")
82
83         return TransformerEncoder(embed_dim=embed_dim,
84                                   num_heads=self.num_heads,
85                                   layers=max(self.layers, layers),
86                                   attn_dropout=attn_dropout,
87                                   relu_dropout=self.relu_dropout,
88                                   res_dropout=self.res_dropout,
89                                   embed_dropout=self.embed_dropout,
90                                   attn_mask=self.attn_mask)
91
92     def forward(self, x_l, x_a, x_v):
93         """
94         text, audio, and vision should have dimension [batch_size, seq_len, n_features]
95         """
96         x_l = F.dropout(x_l.transpose(1, 2), p=self.embed_dropout, training=self.training)
97         x_a = x_a.transpose(1, 2)
98         x_v = x_v.transpose(1, 2)
99
100         # Project the textual/visual/audio features
101         proj_x_l = x_l if self.orig_d_l == self.d_l else self.proj_l(x_l)
102         proj_x_a = x_a if self.orig_d_a == self.d_a else self.proj_a(x_a)
103         proj_x_v = x_v if self.orig_d_v == self.d_v else self.proj_v(x_v)
104         proj_x_a = proj_x_a.permute(2, 0, 1)
105         proj_x_v = proj_x_v.permute(2, 0, 1)
106         proj_x_l = proj_x_l.permute(2, 0, 1)
107
108         if self.lonly:
109             # (V,A) --> L
110             h_l_with_as = self.trans_l_with_a(proj_x_l, proj_x_a, proj_x_a) # Dimension (L, N, d_l)
111             h_l_with_vs = self.trans_l_with_v(proj_x_l, proj_x_v, proj_x_v) # Dimension (L, N, d_l)
112             h_ls = torch.cat([h_l_with_as, h_l_with_vs], dim=2)
113             h_ls = self.trans_l_mem(h_ls)
114             if type(h_ls) == tuple:
115                 h_ls = h_ls[0]
116             last_h_l = last_hs = h_ls[-1] # Take the last output for prediction
117
118         if self.aonly:
119             # (L,V) --> A
120             h_a_with_ls = self.trans_a_with_l(proj_x_a, proj_x_l, proj_x_l)
121             h_a_with_vs = self.trans_a_with_v(proj_x_a, proj_x_v, proj_x_v)
122             h_as = torch.cat([h_a_with_ls, h_a_with_vs], dim=2)
123             h_as = self.trans_a_mem(h_as)
124             if type(h_as) == tuple:
125                 h_as = h_as[0]
126             last_h_a = last_hs = h_as[-1]
127
128         if self.vonly:
129             # (L,A) --> V
130             h_v_with_ls = self.trans_v_with_l(proj_x_v, proj_x_l, proj_x_l)
131             h_v_with_as = self.trans_v_with_a(proj_x_v, proj_x_a, proj_x_a)
132             h_vs = torch.cat([h_v_with_ls, h_v_with_as], dim=2)
133             h_vs = self.trans_v_mem(h_vs)
134             if type(h_vs) == tuple:
135                 h_vs = h_vs[0]
136             last_h_v = last_hs = h_vs[-1]
137
138         if self.partial_mode == 3:
139             last_hs = torch.cat([last_h_l, last_h_a, last_h_v], dim=1)
140
141         # A residual block
142         last_hs_proj = self.proj2(F.dropout(F.relu(self.proj1(last_hs)), p=self.out_dropout, training=self.training))
143         last_hs_proj += last_hs
144
145         output = self.out_layer(last_hs_proj)
146         return output, last_hs
```

